AI and the Self: Human Identity, Authenticity, and Agency in the Age of AI

Abstract

Tagline: How does AI transform the perception, construction, extension, destabilization, and representation of the self?

1 Workshop Topic, Motivation, and Relevance

Large language models, recommender systems, and multimodal agents are no longer merely external tools. Rather, they are becoming interlocutors, cognitive extensions [3], mirrors, collaborators, and social actors [2] that can mediate personal awareness, preferences, behavioral patterns, and the self in general, a concept some authors have referred to as the "filter bubble" [7], the "relational self" [9], and the "algorithmic self" [5; 11; 7]. With increasingly capable AI systems, understanding how AI constructs, extends, and transforms the self has become necessary to preserve individuality, agency, authenticity, and human flourishing.

The scholarly conversation on AI and selfhood remains largely confined to normative ethics and philosophical inquiry [5; 8]. What is missing is an interdisciplinary engagement that brings empirical rigor to questions that are currently treated as purely conceptual. Specifically, the field lacks the computational infrastructure that would make these questions tractable: formal models of how AI systems alter preference formation and belief revision, datasets capturing longitudinal changes in user identity and decision-making across AI-mediated interactions, benchmarks for assessing the degree to which systems like recommenders or conversational agents reshape rather than reflect user values, and evaluation protocols capable of detecting subtle shifts in memory, self-narrative, and agentive capacity [10; 2]. Building this infrastructure requires sustained dialogue between AI researchers, cognitive scientists, psychologists, sociologists, and ethicists, not as an afterthought, but as a constitutive feature of how the field advances.

To the best of our knowledge, no NeurIPS workshop directly theorizes how AI systems construct, extend, or destabilize the human self. In [12], presented at the NeurIPS Socially Responsible Language Modelling Research (SoLaR) workshop (2024), Ward take an important step in this direction by examining how LLMs may influence users beliefs and behavioral dispositions through repeated interaction, highlighting risks of subtle value steering and preference shift [1; 4]. However, even this line of work remains primarily focused on model behavior and downstream social impact, rather than developing a unified computational theory of self-modification dynamics. More broadly, recent workshop ecosystems across NeurIPS, ICML, ICLR, and ACL increasingly recognize issues of safety, fairness, and human-centered design, but they still lack a formal framework that treats the self as an evolving state variable shaped by iterative AI interaction.

Building on the interdisciplinary workshop *AI and the Self: Interdisciplinary and Cross-Cultural Perspectives* we organized at the University of Bonn in June June 2026 [1], this proposal aims to translate its conceptual outcomes[2] into a concrete research agenda for the broader machine learning community. Rather than treating the workshops findings as purely interpretive or thematic, we position them as design constraints for a new class of ML systems that explicitly account for their role in shaping human self-models. This shift enables a tighter coupling between empirical observations of AI-mediated selfhood and the methodological toolkits of modern machine learning, including representation learning, sequential decision-making, and human-in-the-loop optimization.

[1] https://www.cst.uni-bonn.de/en/events/workshop-ai-and-the-self

[2] Ten areas in which AI influences the self: AI as a mirror; cognitive extension and skill change; identity formation and loss; data selves, avatars, and digital embodiment; emotion and dependence; knowledge, expertise, and human uniqueness; culture, language, and worldview; reality, consciousness, and perception; ethics, inequality, and professional authorship; and philosophical, metaphysical, and spiritual conceptions of self

Preprint. Under review.

More specifically, we propose to bring these insights into NeurIPS-facing research through the lens of self-dynamics as a learnable and measurable object, where interaction histories are not only signals for preference inference but also trajectories of evolving identity and agency. This reframing opens the door to studying when and how optimization processes in recommender systems, language models, and agents induce latent drift in user self-representation and how such effects can be systematically characterized, predicted, or constrained. In this way, the workshop acts as a conceptual bridge toward formalizing selfhood as an emergent property of coupled humanAI learning systems, enabling new directions at the intersection of sequence modeling, distribution shift, and interactive learning.

2 Workshop Agenda and Structure

The workshop will organize contributions around five linked themes connecting technical AI research with empirical, social, cultural, and philosophical inquiry.

(i) Measurement and evaluation. How can we measure AI's effects on self-concept, agency, confidence, dependence, identity formation, and skill change, and evaluate whether systems support reflection and autonomy rather than manipulation or over-reliance?

(ii) Technical design and interaction. How should AI systems implement memory, personalization, uncertainty, abstention, refusal, and "not knowing" while avoiding over-agreement, unhealthy attachment, deceptive empathy, and culturally narrow representations of the self?

(iii) Culture, language, and worldview. How do AI systems encode assumptions about the human self, intelligence, emotion, and agency, and how can they better support diverse human identities by avoiding cultural, moral, and linguistic flattening [6]?

(iv) Professional, educational, and epistemic transformation. How does AI change learning, authorship, expertise, creativity, classroom culture, professional identity, and the meaning of human competence when it acts as a co-author, tutor, collaborator, or co-worker?

(v) Philosophical, ethical, legal, and spiritual foundations. Does AI challenge human uniqueness, autonomy, responsibility, consciousness, authorship, or personhood, and how should ethical or legal responsibility be assigned when self-expression and knowledge production are co-created with AI?

We will invite **short and long non-archival papers** from technical, empirical, socio-cultural, philosophical, and application (in education, mental and physical health, and so on) oriented perspectives. Submissions may include technical work on model design and evaluation; empirical studies of AI-mediated reflection, development, education, and dependence; conceptual papers on selfhood, agency, authorship, responsibility, and culture; and demos of reflective AI systems, avatars, and cultural-preservation tools. We especially welcome work-in-progress, early findings, new problem framings, and interdisciplinary provocations that refine the AI and Self research agenda. Submissions must be unpublished and not under review elsewhere. Work accepted to the main conference may not be presented here, including within invited talks, and the call will state this explicitly.

Tentative Schedule. Below is a tentative schedule of the one-day workshop:

- **(09:00–09:15) Opening** remarks and workshop agenda.
- **(09:15–10:15) Keynote 1: The Philosophical Self in the Age of AI.** Tentative speakers: Amanda Askell or David Chalmers or Thomas Metzinger.
- **(10:15–11:15) Panel: Interdisciplinary perspectives on the experiential self.** Moderated discussion on qualitative, phenomenological, philosophical, and humanistic views on AI-mediated selfhood. Potential panel: Markus Gabriel, Thomas Metzinger, Munmun De Choudhury, Anders Søgaard, Sheffali Gulati.
- **(11:15–12:45)** Coffee break and **Poster Session I**.
- **(12:45–14:15)** Lunch.
- **(14:15–15:00) Keynote 2: The Empirical Self in AI Systems.** Tentative speakers: Melanie Mitchell or Munmun De Choudhury.
- **(15:00–16:30) Spotlight Prentations** of the accepted papers (upto 9, 10 mins each).
- **(16:30–18:00)** Coffee break and **Poster Session II**.
- **(18:00–18:30) Closing:** Best paper awards and next steps for the AI x Self community.

3 Invited Speakers

We invited an interdisciplinary group of 8 speakers for in-person talks, spanning AI, Eastern and Western philosophy, neuroscience, cognitive science, and social science (2 already confirmed).[3]

1. Amanda Askell (https://askell.io, **Tentative**) is a philosopher working on finetuning and AI alignment at Anthropic. Her work focuses on shaping model behavior, and human-AI interaction.

2. David Chalmers (https://consc.net, **Tentative**) is a Professor of Philosophy and Neural Science at New York University, recently focusing on studying AI consciousness, welfare, and digital selves.

3. Munmun De Choudhury (http://www.munmund.net, **Confirmed**) is a Professor at Georgia Tech, focusing on computational social science in mental well-being, self-reflection, and AI reliance.

4. Markus Gabriel (https://markus-gabriel.com, **Confirmed**) is the Director of the Center for Science and Thought, University of Bonn, focusing on selfhood, agency, and human-AI relations.

5. Sheffali Gulati (https://www.ilae.org/sheffali-gulati, **Tentative**) is the Chief of the Child Neurology Division at the All India Institute of Medical Sciences, focusing on AI, neurodevelopment, and child wellbeing.

6. Thomas Metzinger (https://thomasmetzinger.com, **Tentative**) is a Professor Emeritus at Johannes Gutenberg University in Mainz, who focuses on the philosophy of mind, and applied ethics.

7. Melanie Mitchell (https://melaniemitchell.me, **Tentative**) is a Professor at the Santa Fe Institute bridging AI, and cognitive science, to evaluate LLM metacognition, and skills in analogy-making.

10. Anders Søgaard (https://anderssoegaard.github.io, **Tentative**) is Professor of NLP and ML at University of Copenhagen, recently focusing on fairness, trust, and societal impact of AI personas.

4 Other Details

Attendance. We estimate ~180 attendees from AI, HCI, Cognitive Science, Philosophy, and more.
Other Proposals. None of the authors have submitted any other workshop proposals to NeurIPS 2026, or related proposals to any other venue.
Timeline. Call for papers: 12 July, Submission deadline: 29 August (AoE). Reviewing: 1-21 September. Author notification: 28 September. Camera-ready and final schedule: October 2026.
Location preference in descending order of preference: *Sydney*, *Paris*.
Technical needs. The workshop requires standard setup: projector, screen, microphones, Wi-Fi, and a room with poster boards and space. Invited talks may be recorded if permitted.
Why now. AI systems are becoming intimate and personalised faster than we are building methods to measure their effects on self-concept, agency, and dependence. Without a shared agenda now, the field risks fragmenting and shipping systems that reshape how people understand themselves before we can evaluate or govern those effects.

4.1 Diversity Statement

Diversity is a commitment and a part of the scientific content of this workshop: **(i) Organizers:** represents diversity across geography, institution type, seniority, gender, and cultural background. **(ii) Speakers:** span technical and humanistic views: NLP, HCI, cognitive science, philosophy, cultural studies, and ethics, with a focus on diversity in gender, geography, seniority, institution type, and method. **(iii) Participants:** We welcome early-career researchers, practitioners, and scholars from NLP, HCI, cognitive science, philosophy, education, sociology, design, and healthcare, especially encouraging work on low-resource languages, non-Western and cross-cultural concepts of self, disability and access, youth and development, inequality, emotional dependence, and underrepresented value systems. We will also promote the workshop through inclusive AI communities, such as *Women in ML*, *DisAbility in AI*, etc, and seek support for registration or travel grants.

[3] Invited Speakers and Organizers are listed alphabetically by surname. Listing order does not indicate priority, seniority, contribution, or confirmation status.

5 Organizers

1. Akhil Arora is a professor of computer science at Aarhus University, where he leads the CLAN for AI Research on Language and Networks (or "CLAN" for short). He is a fellow of Copenhagen Center for Social Data Science, an affiliate of ELLIS Unit Copenhagen and Pioneer Centre for AI, and a formal collaborator of the Wikimedia Foundation. Akhil's research lies broadly in human-centered and trustworthy AI, with a focus on structured and multi-agent reasoning, multilinguality, culturally grounded language technologies, and systems for reliable and cost-efficient LLM inference. His work has been published in top-tier AI venues including ICML, ACL, EMNLP, NAACL, and WWW. He has co-organized multiple workshops collocated with top venues including ACL, EMNLP, and SIGMOD. Akhil is a director of the P1-programs on Green AI and AI & Society.

2. Monojit Choudhury is a Professor of Natural Language Processing, and the Chief Scientist of the Institute of Agriculture and AI at Mohamed bin Zayed University of Artificial Intelligence (MBZUAI), Abu Dhabi. Prior to this, he was a Principal Scientist at Microsoft Research Lab and Microsoft Turing, India. His research interests lie in the intersection of NLP, Social and Cultural aspects of Technology use, and Ethics. He is a member of the ACL Ethics Committee. Monojit has organized numerous workshops at *ACL conferences, including the TextGraphs series and Computational Processing of Code-mixing series.

3. Mukund Choudhary is a PhD candidate at MBZUAI researching how LLMs structure language or "LLM linguistics", currently focused on ways to compare LLM and human linguistic and language-learning behaviour and distinguishing it from internal mechanisms. Apart from regularly reviewing and volunteering at *ACL conferences, he has also headed the organization, coordination, and outreach for the Panini Linguistics Olympiad for high-school students in India.

4. Abdoul Jalil Djiberou Mahamadou (he/him) is a Postdoctoral Associate at MBZUAI working on personalization, cultural sensitivity, and inclusion in AI. He has previously worked on the ethics of AI in healthcare, algorithmic bias and fairness, participatory AI, AI and neuroscience, causal inference, and uncertainty modelling. He is an organizing comittee member of the incoming AACL and EMNLP tutorial on ethics in NLP.

5. Jacki O'Neill is the Director and founding Lab Director of Microsoft Research Africa, Nairobi (formerly MARI). She applies ethnographic and HCI methods to build culturally grounded, human-centered AI and translate fieldwork into systems across work, health, and agriculture. She received the ACM SIGCHI Societal Impact Award (2026) and has served on organizing/program committees for CHI, CSCW, ICTD, and ECSCW, including co-organizing the Nairobi workshop behind the *AI and the Future of Work in Africa* white paper.

6. Dr. Haneesha Pinnamaraju is the Medical Director at SHINE Child Development Center, Dubai, and a Developmental & Behavioral Pediatrician. Through her work and expertise on neurodiversity and inclusive education, she contributes developmental perspectives to discussions on the design, use, and societal impact of AI, with a focus on ensuring that AI technologies support inclusion, human flourishing, and the diverse developmental needs of children and adolescents.

7. Sandrine R. Schiller Sandrine R. Schiller is a postdoctoral fellow at the Center for Philosophy of AI at the University of Copenhagen. Her research explores the implications of intimate relationships with AI chatbots and how generative AI transforms the conditions of selfhood and intersubjective relations. Her research also extends to philosophical questions in mechanistic interpretability and the risks of engagement optimization. Her research has been published by AAAI/ACM and NAACL, as well as more traditional venues for philosophy.

6 Program Committee

Table 1 lists[4] 32 tentative Program Committee members. Upon acceptance of this workshop, we will recruit at least 10 additional PhD-level reviewers to ensure that each submission receives at least three reviews.

Reviews will be managed through OpenReview, enabling conflict-of-interest handling and post-workshop discussion of accepted papers. Accepted papers will be made available on the workshop website, with optional links to project pages, code, videos, or other author-provided materials.

[4]Tentative PC Members are listed alphabetically by first name. Listing order does not indicate priority, seniority, or confirmation status.

Table 1: Tentative List of Program Committee Members

Name	Designation	Affiliation
Aikaterini (Katerina) Fotopoulou	Professor	University College London
Alham Aji Fikri	Asst Professor	MBZUAI
Animesh Mukherjee	Professor	IIT Kharagpur
Ankita Maity	PhD student	Aarhus University
Anna Helene	Postdoctoral Researcher	Copenhagen Center for Social Data Science
Arnav Arora	PhD student	University of Copenhagen & Apple
Bert Heinrichs	Professor	University of Bonn
Camilo Miguel Signorelli	Postdoctoral Researcher	University of Copehagen
Chelsea Haramia	Senior Research Fellow	University of Bonn
Chris Frith	Professor	University of London
Dustin Wright	Assistant Professor	Aalborg University
Hanne De Jaegher	Visiting Fellow	University of Sussex
Iwan Williams	Postdoctoral Researcher	Centre for Philosophy of AI, Copenhagen
Julia Kreutzer	Research Scientist	Cohere for AI
Julian Kiverstein	Assistant Professor	University of Amsterdam
Jurgita Imbrasaite	Senior Research Fellow	University of Bonn
Karl Friston	Professor	University College London
Kenneth Enevoldsen	Assistant Professor	Aarhus University
Lizhen Qu	Senior Lecturer	Monash University
Laurits Lyngbæk	PhD student	Aarhus University
Maria Antoniak	Assistant Professor	University of Colorado Boulder
Marjolein Lanzing	Assistant Professor	University of Amsterdam
Mina Almasi	PhD Student	Aarhus University
Murray Shanahan	Principal Scientist	Google DeepMind
Niket Tandon	Principal Researcher	Microsoft Research
Prashant Kumar	Junior Fellow	Dept of Philosophy, University of Bonn
Siddhesh Pawar	PhD Student	University of Copenhagen
Sarah Masud	Postdoctoral Associate	University of Copenhagen
Sougata Saha	Postdoctoral Associate	MBZUAI
Tianyi Hu	PhD student	Aarhus University
Zeerak Talat	Assistant Professor	University of Edinburgh
Zoe Horlacher	PhD student	Copenhagen Center for Social Data Science

Each submission will receive at least three double-blind reviews, and no reviewer will be assigned more than three papers. Acceptance and awards will be decided collectively by the organizers. Any organizer with a conflict of interest, including authorship or institutional affiliation, will be excluded from decisions involving the relevant paper.

References

[1] Emily M Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*, pages 610–623, 2021.

[2] Monojit Choudhury. Generative ai has a language problem. *Nature human behaviour*, 7(11):1802–1803, 2023.

[3] Andy Clark and David Chalmers. The extended mind. *analysis*, 58(1):7–19, 1998.

[4] Munmun De Choudhury, Sachin R Pendse, and Neha Kumar. Benefits and harms of large language models in digital mental health. *arXiv preprint arXiv:2311.14693*, 2023.

[5] Jeena Joseph. The algorithmic self: How AI is reshaping human identity, introspection, and agency. *Frontiers in Psychology*, Volume 16 - 2025, 2025. ISSN 1664-1078. doi: 10.3389/fpsyg.2025.1645795. URL https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2025.1645795.

[6] Pratik Joshi, Sebastin Santy, Amar Budhiraja, Kalika Bali, and Monojit Choudhury. The state and fate of linguistic diversity and inclusion in the nlp world. In *Proceedings of the 58th annual meeting of the association for computational linguistics*, pages 6282–6293, 2020.

[7] Eli Pariser. *The filter bubble: What the Internet is hiding from you*. penguin UK, 2011.

[8] S Santhya, K Sashtiga, and N Senthil Kumar. The evolution of selfhood in the digital era: Philosophical reflections on identity, authenticity and technology-mediated human experience. *International Journal of Research and Scientific Innovation (IJRSI)*, 12(9), 2025.

[9] Sandrine R Schiller, Camilo Miguel Signorelli, and Filippos Stamatiou. The intercepted self: How generative ai challenges the dynamics of the relational self. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, volume 8, pages 2284–2291, 2025.

[10] Inhwa Song, Sachin R Pendse, Neha Kumar, and Munmun De Choudhury. The typing cure: Experiences with large language model chatbots for mental health support. *Proceedings of the ACM on Human-Computer Interaction*, 9(7):1–29, 2025.

[11] Grace Leonora Turtle, Elisa Giaccardi, and Roy Bendor. Undoing gracia: Queering the self in the algorithmic borderlands. *AI & SOCIETY*, 41(4):3649–3666, 2026.

[12] Francis Rhys Ward. Towards a theory of ai personhood. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 27680–27688, 2025.